

Tshlab

Analysis Report

Purpose

To become more familiar with process control and signalling by finishing a tiny linux shell.

Work--complete functions below:

- `eval()` : main routine that parses and interprets the cmdline
- `sigchld_handler`, `sigint_handler`, `sigstp_handler`

Requirements:

- redirection of input & output
- pipes not needed
- built-commands & job control
 - `quit`
 - `jobs`
 - `bg/fg` job
- *reap all of its zombie children*

File Hierachy

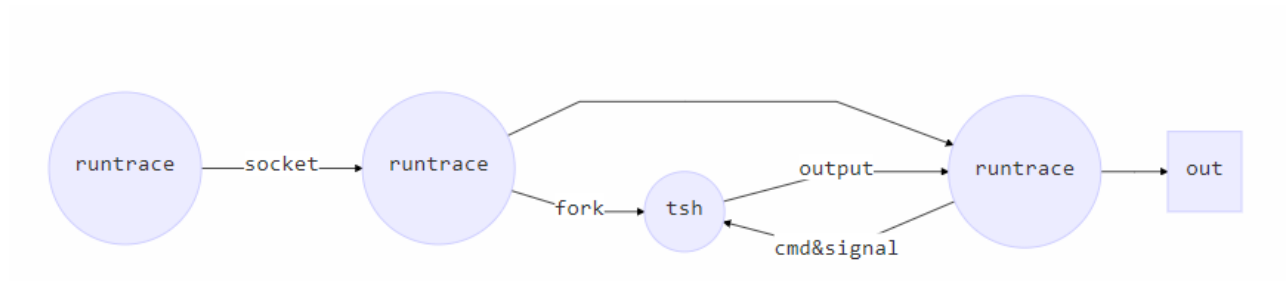
tshlab/src/	Main source codes, some of which will be copied to handout and sent to students
tshlab/writeup/	Writeup of Lab sent to students
tshlab/patches/	Patches added by PKU // to fix bugs ?
tshlab/test/	Test
tshlab/test-autograder/	Test
tshlab/online-solutions/	Online solutions // may be used for cheat checker ?
tshlab/handin/	Directory of handins from students
tshlab/handout/	Directory sent to student
tshlab-rubric.txt	Other code rules
tshlab/Makefiles and other config files for autolab	

Main Part -- tshlab/src

This directory contains the source for the driver program to test the tiny shell.

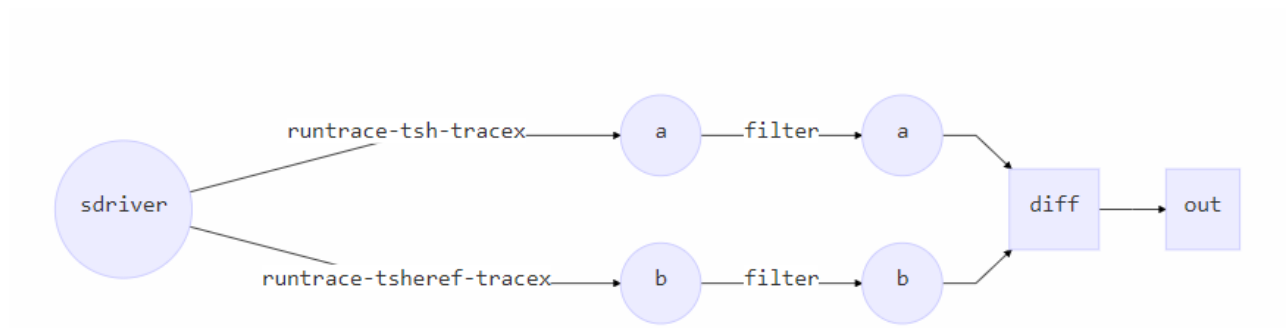
- *runtrace.c* is the program used to run and test *tsh*. It forks a child to execve *tsh*, and the parent process uses socket pairs to send signal and data to child process(*tsh*). The *datafd* socket pair is used to send and receive datas, in other words, the input and output of *tsh*. The *syncfd* socket pair is used to send and receive synchronize signals(i.e. WAIT SIGNAL).

This synchronization can prevent from race condition. For example, parent process sends `"/bin/myspin1 10"` to *tsh* through *datafd* and then hope to send a SIGINT to it. But when the SIGINT arrives, *tsh* may not have started *myspin1*, then the *myspin1* never stops and time out. So we send a SIGNAL from *myspin1* and call WAIT in parent process before we send SIGINT.



- *sdriver.c* is the main routine that runs the driver.

The driver tests the functionality of a shell by calling *runtrace* for multiple times. It uses trace files as input, and compare the output of *tsh* with that from the *tshref*, ignoring blank spaces and PID. We consider the *tsh* right on some trace if and only if the output matches.



- *trace*.c* are trace files used to test the tiny shell. The details of them are discussed in writeup.
- *my*.c* are helper programs used by traces.
- *fork.c* is the wrapper for *fork* that introduces non-determinism in the order that parent and child are executed.
- *elide.pl* is used to elide code between `"$begin <handout>"` and `"$end <handout>"` to generate *tsh-handout.c* for students.
- *tsh-broken.c* is a broken version of *tsh* which uses `'tcsetpgrp'`.

Upgrade

fix bug

- In function *builtin_cmd()*, opened file descriptor isn't closed.

tshlab/src/tsh.c

```

11 (void) flush();
12 out_fd = open(tok->outfile, O_WRONLY | O_CREAT | O_TRUNC, 0644);
13 if(out_fd == -1) {
14     (void) fprintf(stderr, "Error: %s Couldn't open file\n",
15                     tok->outfile);
16     return 1;
17 }
18 }
19
20 switch (tok->builtins) {
21 case BUILTIN_QUIT:
22     fflush(stdout);
23     fflush(stderr);
24     exit(0);
25 case BUILTIN_JOBS:
26     listjobs(job_list, out_fd);
27     return 1;
28 case BUILTIN_FG:
29 case BUILTIN_BG:
30     do_bgfg(tok->argv, out_fd);
31     return 1;
32 default:
33     /* not a builtin command */
34     (*input_fd) = in_fd;
35     (*output_fd) = out_fd;
36     return 0;
37 }
38
39 return 0; /* not a builtin command */

```

- Gcc error -- sprintf buffer isn't big enough.

```

lzzz@DESKTOP-5DIRHK6:/mnt/c/Users/Joe/OneDrive - pku.edu.cn/Grade_3/OSP/OSP/Project/tshlab-handout$ make
/usr/bin/gcc -Wall -g -Werror -Wl,--wrap,fork -o tsh tsh.c fork.c
tsh.c: In function 'listjobs':
tsh.c:713:31: error: 'sprintf' may write a terminating nul past the end of the destination [-Werror=format-overflow=]
    sprintf(buf, "%s\n", job_list[i].cmdline);
                                ^
tsh.c:713:13: note: 'sprintf' output between 2 and 1025 bytes into a destination of size 1024
    sprintf(buf, "%s\n", job_list[i].cmdline);
    ^~~~~~
cc1: all warnings being treated as errors
Makefile:19: recipe for target 'tsh' failed
make: *** [tsh] Error 1

```

- When job numbers exceed MAXJOBS, segmentation fault happens.

```

tshlab/src/tsh.c MODIFIED
Side-by-side diff View file Comment ...

326 326     signals handlers can run again. Notice that the global job
327 327     list is protected because the signals for the three
328 328     handlers that manipulate it are blocked. */
329 329     addjob(job_list, pipe_counts + 1, pids, status, cmdline);
330 330
331 331     if (!addjob(job_list, pipe_counts + 1, pids, status, cmdline)) {
332 332         return;
333 333     }
334 334     char buf[MAXLINE];
335 335     struct job_t *j = getjobpid(job_list, pids[0]);
336 336     if (!bg) {

```

- make "pkill" only kill its process.

...

88

99

1010

1111

1212

1313

1414

1515

1616

...

tshlab/src/trace10.txt

MODIFIED

WAIT

-/bin/echo -e tsh\076 /usr/bin/pkill myspin1

+ /bin/echo -e tsh\076 /usr/bin/pkill -u \$USER myspin1

NEXT

-/usr/bin/pkill myspin1

+ /usr/bin/pkill -u \$USER myspin1

NEXT

SIGNAL

Additionally, we use SIGUSR1 to synchronize and prevent from race condition.

```

    }
    Signal(SIGTTIN, SIG_DFL);
    Signal(SIGTTOU, SIG_DFL);
-    if (i == 0)
+    if (i == 0) {
        setpgid(0, 0);
        /* send a SIGUSR1 to parent */
        kill(getppid(), SIGUSR1);
    }
    else {
        setpgid(0, childgid);
    }
}

```

For simplicity, *tsh* need not to support built-in commands with pipes.

Job struct & perl

Modify job struct and related operations(e.g. printf).

<pre> 61 struct job_t { /* The job struct */ 62 - pid_t pid; /* job PID */ 63 int jid; /* job ID [1, 2, ...] */ 64 int state; /* UNDEF, BG, FG, or ST */ 65 char cmdline[MAXLINE]; /* command line */ 66 }; 67 struct job_t job_list[MAXJOBS]; /* The job List */ </pre>	<pre> 61 struct job_t { /* The job struct */ 62 + int pgid; /* The number of processes in a job */ 63 + int procnum; /* job ID [1, 2, ...] */ 64 int jid; /* job ID [1, 2, ...] */ 65 int state; /* UNDEF, BG, FG, or ST */ 66 + pid_t pids[MAXPIPES + 1]; /* job PID */ 67 char cmdline[MAXLINE]; /* command line */ 68 }; 69 struct job_t job_list[MAXJOBS]; /* The job List */ </pre>
---	--

Modify sdriver to filter "(pid, pid, ..., pid)".

```

* These transformations allow us to do diffs on the outputs of
* different runs of different shells.
*/
-#define PERLPROG "while(<>){chomp; s/\\s+//g; s/\\(\\d+\\)/\\(PID\\)/g; print \"$_\\n\"}
+#define PERLPROG "while(<>){chomp; s/\\s+//g; s/\\(\\.\\*\\)/[PIDS]/g; print \"$_\\n\"}

/*****
* Global variables

```

trace

fix trace10:

```

SIGINT
NEXT

-/bin/echo -e tsh\076 /bin/sh -c \047/bin/ps h | /bin/fgrep -v grep | /bin/fgrep mysplit\047
+/bin/echo -e tsh\076 /bin/sh -c \047/bin/ps h \0174 /bin/fgrep -v grep \0174 /bin/fgrep mysplit\047
NEXT
/bin/sh -c '/bin/ps h | /bin/fgrep -v grep | /bin/fgrep mysplit'
NEXT

```




















Add more traces to test for pipes:

- *trace25* Pipe - use commands with a pipe to xargs
- *trace26* Pipe - test if children are forked by tsh
- *trace27* Pipe - background jobs with pipe
- *trace28* Pipe - use commands with multiple pipes

There are four traces which test for pipes at 4 pts each. Only one trace tests for multiple pipes.

Source Code & Git Repositroy & Writeup

See git repository.

Author	Commit	Message
 Lzzz	b24eb1e	modify writeup -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	cfeea76	add comment by Cheng Sheng <1600012909@pku.edu.cn>
 Lzzz	ff33525	modify config.h, tsh head -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	1c98b8e	add environment -- implemented by Cheng Sheng <1600012909@pku.edu.cn>
 Lzzz	75c92dd	modify PERLPROG to ignore anything between "()" -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	36f6982	fix fgpid bug, add fgpid in job struct -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	bc6fc85	fix race bug -- add SIGUSR1 to synchronize -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	302b7e9	fix "loop initial declaration" bug, robust code fix -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	7a11f67	fix pkill bug -- should implement ENV \$USER -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	0f39e50	add traces to test pipes -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	0530e85	fix "too many jobs" segmentation fault bug -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	da5722d	fix code to more robust -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	6898cc3	fully implement pipes & add helper functions -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	6f7d300	implement pipes -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	192e465	fix traces: substitute with \0174 -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	21679d5	fix bug: tsh.c -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	164313b	rewrite tsh.c and then add deal_pipe
 Lzzz	2484bcf	fix bug: builtin_cmd may not close file descriptors -- SoB Li Zhuo <lizhmq@pku.edu.cn>
 Lzzz	294f118	init

To-do List

- Builtin-cmd & pipes

For simplicity, current tsh doesn't support builtin-cmd with pipes.

- Multi-process control

For foreground job, tsh will wait for all processes to terminate. In other words, a job is considered terminated if all of its processes has terminated.

But what's the status of a job if some processes in it are stopped while others are running? This is undefined in tsh.

- Background job read/write

Jobs in tsh is actually back ground job in bash(tsh is the foreground job), so they cann't read/write from/to stdin/stdout. Otherwise SIGTTIN/SIGTTOUT will be raised.

- Evaluation

We judge the tsh by its output which may be improved by other methods.

