

# Lab06 Debug

李拙

# 1

- 用print\_hex\_dump打印 boot\_command\_line 和 saved\_command\_line 的地址空间内容

# 1.1 print\_hex\_dump

- 将二进制内容以16进制打印，并可选打印ascii字符
- 主要参数：
  - `const char *level, const char *prefix_str`
  - `const void *buf, size_t len, bool ascii`

# 1.1 print\_hex\_dump

\* E.g.:

```
* print_hex_dump(KERN_DEBUG, "raw data: ", DUMP_PREFIX_ADDRESS,  
*               16, 1, frame->data, frame->len, 1);
```

\*

\* Example output using %DUMP\_PREFIX\_OFFSET and 1-byte mode:

```
* 0009ab42: 40 41 42 43 44 45 46 47 48 49 4a 4b 4c 4d 4e 4f  @ABCDEFGHIJKLMNO
```

## 1.2 command\_line

- `cmdline`是内核启动的参数配置字符串
- 格式为“value1, param=value2”
- 在`start_kernel()`中处理
- 定义在`init/main.c`

## 1.2 command\_line

- `char __initdata boot_command_line[COMMAND_LINE_SIZE];`
- `char *saved_command_line;`
- `bootcmdline`定义时加上宏`__initdata`, 链接时会放在`.init.data`段, `init`阶段之后释放
- `savedcmdline`作用是保存未改动的`cmdline`

## 1.3 导出cmdline

- bootcmdline 和 savedcmdline都没有导出
- 修改源码:
  - EXPORT\_SYMBOL(boot\_command\_line);
  - EXPORT\_SYMBOL(saved\_command\_line);
- 重新编译、替换内核
- (导出的符号处于内核态)

## 1.4 编写模块

- 不知道要打印的长度
- 自己写了一个strlen...

```
int my_strlen(char *s)
{
    int i = 0;
    while (s[i++]);
    return i - 1;
}
```



## 1.4 编写模块

- 注: print\_hex\_dump要 #include <linux/printk.h>

```
static int __init dump_init(void)
{
    print_hex_dump(KERN_INFO, "boot_command_line: ", DUMP_PREFIX_NONE, 16, 1,
boot_command_line, my_strlen(boot_command_line), 1);
    print_hex_dump(KERN_INFO, "saved_command_line: ", DUMP_PREFIX_NONE, 16, 1,
saved_command_line, my_strlen(saved_command_line), 1);
    return 0;
}
```

# 1.5 result

```
[ 1554.618976] boot_command_line: 6c 74 5f 74 61 72 67 65 74  
lt_target  
[ 1554.619689] saved_command_line: 42 4f 4f 54 5f 49 4d 41 47 45 3d 2f 62 6f 6f 74  
BOOT_IMAGE=/boot  
[ 1554.620258] saved_command_line: 2f 76 6d 6c 69 6e 75 7a 2d 34 2e 31 39 2e 30 2b  
/vmlinuz-4.19.0+  
[ 1554.620828] saved_command_line: 20 72 6f 6f 74 3d 55 55 49 44 3d 63 64 37 62 35  
root=UUID=cd7b5  
[ 1554.621371] saved_command_line: 37 65 36 2d 63 35 38 37 2d 31 31 65 38 2d 39 66  
7e6-c587-11e8-9f  
[ 1554.621925] saved_command_line: 31 63 2d 30 38 30 30 32 37 39 61 38 66 66 61 20 1c-  
0800279a8ffa  
[ 1554.622568] saved_command_line: 72 6f 20 6d 61 79 62 65 2d 75 62 69 71 75 69 74 ro  
maybe-ubiquit  
[ 1554.623151] saved_command_line: 79 y
```

## 2 SysRQ

- SysRQ is a 'magical' key combo you can hit which the kernel will respond to regardless of whatever else it is doing, unless it is completely locked up.
- <https://www.mjmwired.net/kernel/Documentation/sysrq.txt>

## 2.1 开启SysRQ

- 编译内核时需设置 `CONFIG_MAGIC_SYSRQ=Y`
- 启动内核后, 修改 `/proc/sys/kernel/sysrq`

```
0 - disable sysrq completely
1 - enable all functions of sysrq
>1 - bitmask of allowed sysrq functions (see below for detailed function
    description):
    2 = 0x2 - enable control of console logging level
    4 = 0x4 - enable control of keyboard (SAK, unraw)
    8 = 0x8 - enable debugging dumps of processes etc.
   16 = 0x10 - enable sync command
   32 = 0x20 - enable remount read-only
   64 = 0x40 - enable signalling of processes (term, kill, oom-kill)
  128 = 0x80 - allow reboot/poweroff
  256 = 0x100 - allow nicing of all RT tasks
```

## 2.2 use SysRq key

- x86: ALT+PrtSc+<command key>
- test: ALT+PrtSc+p – SysRq: Show Regs

## 2.2 use SysRq key

```
2806.675327] sysrq: SysRq : Show Regs
[ 2806.676018] CPU: 2 PID: 0 Comm: swapper/2 Tainted: P      C OE      4.19.0+ #4
[ 2806.676020] Hardware name: innotek GmbH VirtualBox/VirtualBox, BIOS VirtualBox 12/01/2006
[ 2806.676028] RIP: 0010:native_safe_halt+0x6/0x10
[ 2806.676033] Code: 71 ff ff ff 7f 5d c3 65 48 8b 04 25 00 5c 01 00 f0 80 48 02 20 48 8b 00 a8 08 74 8b eb c1 90 90 90 90 90 90 55 48 8
9 e5 fb f4 <5d> c3 0f 1f 84 00 00 00 00 00 55 48 89 e5 f4 5d c3 90 90 90 90 90
[ 2806.676035] RSP: 0018:ffffac9a0069be70 EFLAGS: 00000246 ORIG_RAX: ffffffff8b000000
[ 2806.676038] RAX: ffffffff8ea14620 RBX: 0000000000000002 RCX: 0000000000000001
[ 2806.676040] RDX: ffff8e1f90d23940 RSI: 0000000000000007 RDI: 0000000000000002
[ 2806.676042] RBP: ffffac9a0069be70 R08: 00000000000000ca R09: 0000000000000020
[ 2806.676043] R10: 0000000000000004 R11: 0000000000000000 R12: 0000000000000002
[ 2806.676044] R13: 0000000000000000 R14: 0000000000000000 R15: 0000000000000000
[ 2806.676054] FS: 0000000000000000(0000) GS:ffff8e1f90d00000(0000) knlGS:0000000000000000
[ 2806.676056] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
[ 2806.676057] CR2: 00007f59dbfb21d0 CR3: 00000000a9348000 CR4: 00000000000406e0
[ 2806.676062] Call Trace:
[ 2806.676071]  default_idle+0x22/0x150
[ 2806.676076]  arch_cpu_idle+0x15/0x20
[ 2806.676078]  default_idle_call+0x23/0x30
[ 2806.676083]  do_idle+0x1cb/0x280
[ 2806.676086]  cpu_startup_entry+0x1d/0x20
[ 2806.676090]  start_secondary+0x1ab/0x200
[ 2806.676095]  secondary_startup_64+0xa4/0xb0
[ 2806.676098] CPU#2: active: 0000000000000000
[ 2806.676107] CPU#2:  gen-PMC0 ctrl: 0000000000000000
[ 2806.676108] CPU#2:  gen-PMC0 count: 0000000000000000
[ 2806.676110] CPU#2:  gen-PMC0 left: 0000000000000000
[ 2806.676116] CPU#2:  gen-PMC1 ctrl: 0000000000000000
[ 2806.676117] CPU#2:  gen-PMC1 count: 0000000000000000
[ 2806.676118] CPU#2:  gen-PMC1 left: 0000000000000000
[ 2806.676124] CPU#2:  gen-PMC2 ctrl: 0000000000000000
[ 2806.676126] CPU#2:  gen-PMC2 count: 0000000000000000
[ 2806.676127] CPU#2:  gen-PMC2 left: 0000000000000000
[ 2806.676133] CPU#2:  gen-PMC3 ctrl: 0000000000000000
[ 2806.676134] CPU#2:  gen-PMC3 count: 0000000000000000
[ 2806.676135] CPU#2:  gen-PMC3 left: 0000000000000000
```

### 3 dump\_stack

- 打印内核调用栈
- 从栈顶开始查找，判断地址是否在内核代码段

```
for (; stack < top; stack += 4) {  
    addr = *(unsigned int *)stack;  
    if (kernel_text_addressp(addr))  
        printk(" [<%p>] %pS\n", (void *)addr, (void *)addr);  
}
```

## 3.1 1中的call\_stack

- 在dump\_init()中调用call\_stack();

```
static int __init dump_init(void)
{
    print_hex_dump(KERN_INFO, "boot_command_line: ", DUMP_PREFIX_NONE, 16, 1,
boot_command_line, my_strlen(boot_command_line), 1);
    print_hex_dump(KERN_INFO, "saved_command_line: ", DUMP_PREFIX_NONE, 16, 1,
saved_command_line, my_strlen(saved_command_line), 1);
    return 0;
}
```



## 3.1 1中的call\_stack

- 在dump\_init()中调用call\_stack();

```
[ 3494.130061] Call Trace:
[ 3494.130069]  dump_stack+0x63/0x85
[ 3494.130071]  ? 0xfffffffffc072e000
[ 3494.130074]  dump_init+0xe/0x1000 [dump]
[ 3494.130077]  do_one_initcall+0x4a/0x1c9
[ 3494.130080]  ? _cond_resched+0x19/0x40
[ 3494.130083]  ? kmem_cache_alloc_trace+0x15c/0x1d0
[ 3494.130086]  do_init_module+0x5f/0x206
[ 3494.130088]  load_module+0x2213/0x2b00
[ 3494.130091]  __do_sys_finit_module+0xfc/0x120
[ 3494.130092]  ? __do_sys_finit_module+0xfc/0x120
[ 3494.130094]  __x64_sys_finit_module+0x1a/0x20
[ 3494.130095]  do_syscall_64+0x5a/0x120
[ 3494.130097]  entry_SYSCALL_64_after_hwframe+0x44/0xa9
```

# 4 syslog

- Syslog is a standard for message logging.
- It allows separation of the software that generates messages.
- See linux manual:
  - <http://man7.org/linux/man-pages/man3/syslog.3.html>

## 4.1 截获系统调用并计数

- 按照Lab05的方式截获系统调用
- 修改源码导出sys\_call\_table
- 重新编译、加载内核

```
unsigned int mkdir_count;

asm linkage int our_sys_mkdir(const char *filename, int flags, int mode)
{
    ++mkdir_count;
    return original_call(filename, flags, mode);
}
```

- 用户程序如何获取该变量?

## 4.2 内核空间与用户空间

- 编写系统调用
- 编写驱动程序
- 使用proc文件系统
- 使用虚拟文件系统

## 4.3 修改系统调用

- 与劫持系统调用的方法相同
- 加入内核模块，修改系统调用表

## 4.3 修改系统调用

```
asmlinkage int print_mkdir_count(void)
{
    return mkdir_count;
}
int init_module()
{
    cr0 = read_cr0();
    write_cr0(cr0 & ~CR0_WP);

    original_call = sys_call_table[233];
    sys_call_table[233] = print_mkdir_count;

    write_cr0(cr0);

    return 0;
}
```

## 4.3 修改系统调用

- 劫持系统调用的模块A计数 `EXPORT_SYMBOL(mkdir_count);`
- and
- 用户调用系统调用B返回mkdir\_count
  - `extern unsigned int mkdir_count;`

## 4.4 用户程序

```
1  #include <linux/unistd.h>
2  #include <syscall.h>
3  #include <sys/types.h>
4  #include <stdio.h>
5  #include <unistd.h>
6  #include <syslog.h>
7
8  int main()
9  {
10     unsigned int c0 = syscall(233);
11     unsigned int c1 = c0;
12     unsigned int c2;
13     syslog(LOG_KERN, "Initial count: %d.\n", c0);
14     while (1) {
15         sleep(60);
16         c2 = syscall(233);
17         openlog("mkdir_count_log", LOG_KERN, 0);
18         syslog(LOG_USER | LOG_DEBUG,
19             "%d sys_mkdir calls in last one minute. %d calls in all.\n", c2 - c1, c2 - c0);
20         closelog();
21         c1 = c2;
22     }
23     return 0;
24 }
25
```



## 4.5 查看结果

- `vim /var/log/syslog`

```
0 Oct 29 16:13:49 ubuntu a.out: Initial count: 0.  
1 Oct 29 16:13:49 ubuntu mkdir_count_log: 0 sys_mkdir calls in last one minute. 0  
calls in all.  
2 Oct 29 16:14:00 ubuntu mkdir_count_log: message repeated 66603 times: [ 0  
sys_mkdir calls in last one minute. 0 calls in all.]  
3 Oct 29 16:14:00 ubuntu mkdir_count_log: 1 sys_mkdir calls in last one minute. 1  
calls in all.  
4 Oct 29 16:14:00 ubuntu mkdir_count_log: 0 sys_mkdir calls in last one minute. 1  
calls in all.  
5 Oct 29 16:14:20 ubuntu mkdir_count_log: message repeated 113647 times: [ 0  
sys_mkdir calls in last one minute. 1 calls in all.]  
6 Oct 29 16:14:20 ubuntu mkdir_count_log: 1 sys_mkdir calls in last one minute. 2  
calls in all.  
7 Oct 29 16:14:20 ubuntu mkdir_count_log: 0 sys_mkdir calls in last one minute. 2  
calls in all.  
8 Oct 29 16:14:24 ubuntu mkdir_count_log: message repeated 18025 times: [ 0  
sys_mkdir calls in last one minute. 2 calls in all.]  
9 Oct 29 16:14:24 ubuntu mkdir_count_log: 1 sys_mkdir calls in last one minute. 3  
calls in all.  
10 Oct 29 16:14:24 ubuntu mkdir_count_log: 0 sys_mkdir calls in last one minute. 3  
calls in all.
```

Thanks.