

1

编写模块fp，使用了浮点运算

```
static int __init fp_init(void)
{
    float b = 1.1;
    printk("a: %d\n", a + (int)b);
    return 0;
}
```

硬浮点

编译选项中添加CFLAGS_fp.o=-mhard-float（为避免浮点部分被优化，设置-O0）

```
1  obj-m += fp.o
2  all:
3  ▶      make -C /lib/modules/$(shell uname -r)/build M=$(PWD) CFLAGS_fp.o="-mhard-float -O0" modules
4
5  clean:
6  ▶      make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

反汇编结果：

可以看到汇编中使用了flds, fstps等x87 FPU指令

```
root@ubuntu:~/lab07/fp# objdump -d fp.o
```

```
fp.o:      file format elf64-x86-64
```

```
Disassembly of section .init.text:
```

```
0000000000000000 <init_module>:
 0:  e8 00 00 00 00      callq  5 <init_module+0x5>
 5:  55                  push   %rbp
 6:  48 89 e5            mov     %rsp,%rbp
 9:  48 83 ec 10         sub     $0x10,%rsp
 d:  d9 05 00 00 00 00    flds   0x0(%rip)          # 13 <init_module+0x13>
13:  d9 5d fc            fstps  -0x4(%rbp)
16:  d9 45 fc            flds   -0x4(%rbp)
19:  d9 7d f6            fnstcw -0xa(%rbp)
1c:  0f b7 45 f6         movzwl -0xa(%rbp),%eax
20:  80 cc 0c            or      $0xc,%ah
23:  66 89 45 f4         mov     %ax,-0xc(%rbp)
27:  d9 6d f4            fldcw  -0xc(%rbp)
2a:  db 5d f0            fistpl -0x10(%rbp)
2d:  d9 6d f6            fldcw  -0xa(%rbp)
30:  8b 55 f0            mov     -0x10(%rbp),%edx
33:  8b 05 00 00 00 00    mov     0x0(%rip),%eax    # 39 <init_module+0x39>
39:  01 d0              add     %edx,%eax
3b:  89 c6              mov     %eax,%esi
3d:  48 c7 c7 00 00 00 00 mov     $0x0,%rdi
44:  e8 00 00 00 00      callq  49 <init_module+0x49>
49:  b8 00 00 00 00      mov     $0x0,%eax
4e:  c9                  leaveq %eax
4f:  c3                  retq
```

```
Disassembly of section .exit.text:
```

```
0000000000000000 <cleanup_module>:
 0:  55                  push   %rbp
 1:  48 89 e5            mov     %rsp,%rbp
 4:  90                  nop
 5:  5d                  pop     %rbp
 6:  c3                  retq
```

软浮点

编译选项中添加CFLAGS_fp.o="-msoft-float"

```
1  obj-m += fp.o
2  all:
3  ▶      make -C /lib/modules/$(shell uname -r)/build M=$(PWD) CFLAGS_fp.o="-msoft-float -O0" modules
4
5  clean:
6  ▶      make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

make时有warning: "__fixfsi" undefined

```

root@ubuntu:~/lab07/fp# make
make -C /lib/modules/4.19.0+/build M=/home/lzzz/lab07/fp CFLAGS_fp.o="-msoft-float -O0" modules
make[1]: Entering directory '/home/lzzz/linux'
CC [M] /home/lzzz/lab07/fp/fp.o
Building modules, stage 2.
MODPOST 1 modules
WARNING: "__fixsfsi" [/home/lzzz/lab07/fp/fp.ko] undefined!
LD [M] /home/lzzz/lab07/fp/fp.ko
make[1]: Leaving directory '/home/lzzz/linux'

```

反汇编结果：

原来应该出现浮点调用指令的地方变成了一个函数调用callq，由于没有链接调用地址未知，由warning提示这个函数调用应该是__fixsfsi，功能是将float转换为int。与编写的模块内容相符。

```

root@ubuntu:~/lab07/fp# objdump -d fp.o

fp.o:      file format elf64-x86-64

Disassembly of section .init.text:

0000000000000000 <init_module>:
 0:  e8 00 00 00 00      callq  5 <init_module+0x5>
 5:  55                  push   %rbp
 6:  48 89 e5            mov     %rsp,%rbp
 9:  48 83 ec 08         sub     $0x8,%rsp
 d:  8b 05 00 00 00 00   mov     0x0(%rip),%eax          # 13 <init_module+0x13>
13:  89 45 fc            mov     %eax,-0x4(%rbp)
16:  8b 45 fc            mov     -0x4(%rbp),%eax
19:  50                  push   %rax
1a:  e8 00 00 00 00      callq  1f <init_module+0x1f>
1f:  48 83 c4 08         add     $0x8,%rsp
23:  89 c2              mov     %eax,%edx
25:  8b 05 00 00 00 00   mov     0x0(%rip),%eax          # 2b <init_module+0x2b>
2b:  01 d0              add     %edx,%eax
2d:  89 c6              mov     %eax,%esi
2f:  48 c7 c7 00 00 00 00 mov     $0x0,%rdi
36:  e8 00 00 00 00      callq  3b <init_module+0x3b>
3b:  b8 00 00 00 00      mov     $0x0,%eax
40:  c9                  leaveq
41:  c3                  retq

Disassembly of section .exit.text:

0000000000000000 <cleanup_module>:
 0:  55                  push   %rbp
 1:  48 89 e5            mov     %rsp,%rbp
 4:  90                  nop
 5:  5d                  pop     %rbp
 6:  c3                  retq

```

分析

从上面的结果可以看出，使用硬浮点编译的代码会使用FPU指令，需要机器有支持协处理器，而使用软浮点选项编译的程序会使用glibc的函数来模拟浮点指令，不需要FPU支持。

2

分解DO_INVALID_OP宏

traps.c中定义了DO_ERROR宏来统一建立一些错误信号的错误处理函数，这里将do_invalid_op分离出来。

修改如图：

```
308 #define DO_ERROR(trapnr, signr, str, name) \
309 dotraplinkage void do_##name(struct pt_regs *regs, long error_code) \
310 { \
311     do_error_trap(regs, error_code, str, trapnr, signr); \
312 }
313
+ 314 #define DO_INVALID_OP(trapnr, signr, str) \
+ 315 dotraplinkage void do_invalid_op(struct pt_regs *regs, long error_code) \
+ 316 { \
+ 317     do_error_trap(regs, error_code, str, trapnr, signr); \
+ 318 }
+ 319
320 DO_ERROR(X86_TRAP_DE, SIGFPE, "divide error", divide_error)
321 DO_ERROR(X86_TRAP_OF, SIGSEGV, "overflow", overflow)
~ 322 // DO_ERROR(X86_TRAP_UD, SIGILL, "invalid opcode", invalid_op)
+ 323 DO_INVALID_OP(X86_TRAP_UD, SIGILL, "invalid opcode");
324 DO_ERROR(X86_TRAP_OLD_MF, SIGFPE, "coprocessor segment overrun", coprocessor_segment_overrun)
325 DO_ERROR(X86_TRAP_TS, SIGSEGV, "invalid TSS", invalid_TSS)
326 DO_ERROR(X86_TRAP_NP, SIGBUS, "segment not present", segment_not_present)
327 DO_ERROR(X86_TRAP_SS, SIGBUS, "stack segment", stack_segment)
328 DO_ERROR(X86_TRAP_AC, SIGBUS, "alignment check", alignment_check)
329
```

3

编写普通程序：

```
1  #include <stdio.h>
2
3  int main()
4  {
5      printf("Hello world!\n");
6      return 0;
7  }
```

编译生成a.out，反汇编查看main函数位置之后，在二进制指令中加上一个非法指令FF，再反汇编查看结果：

```
120 000000000000063a <main>:
1  63a: 55          > push    %rbp
2  63b: 48 89 e5    > mov     %rsp,%rbp
3  63e: 48 8d 3d 9f 00 00 00 > lea     0x9f(%rip),%rdi    # 6e4 <_IO_stdin_used+0x4>
4  645: e8 c6 fe ff ff > callq   510 <puts@plt>
5  64a: ff         > (bad)==
6  64b: b8 00 00 00 00 > mov     $0x0,%eax
7  650: c3         > retq    ==
8  651: 66 2e 0f 1f 84 00 00 > nopw    %cs:0x0(%rax,%rax,1)
9  658: 00 00 00 ==
10 65b: 0f 1f 44 00 00 > nopl    0x0(%rax,%rax,1)
```

可以看到其中多出了一条指令：ff (bad)

执行的结果是：

```
Hello world!
Illegal instruction (core dumped)
```

程序在执行printf之后，在ff指令处发生invalid op

修改do_error_trap函数：

在发生SIGILL时打印Never Giveup，由于trap之后回到原来指令，所以将ip加1，跳过这条指令，然后返回。

```
33 static void do_error_trap(struct pt_regs *regs, long error_code, char *str,
32 >         unsigned long trapnr, int signr)
31 {
30 >     siginfo_t info;
29
28 >     RCU_LOCKDEP_WARN(!rcu_is_watching(), "entry code didn't wake RCU");
27
26 >     /*
25 >     * WARN*()s end up here; fix them up before we call the
24 >     * notifier chain.
23 >     */
22 >     if (!user_mode(regs) && fixup_bug(regs, trapnr))
21 >         return;
20
+ 19 >     if (signr == SIGILL) {
+ 18 >         printk("Never Giveup");
+ 17 >         regs->ip += 1;
+ 16 >         return;
+ 15 >     }
+ 14
13 >     if (notify_die(DIE_TRAP, str, regs, error_code, trapnr, signr) !=
12 >         NOTIFY_STOP) {
11 >         cond_local_irq_enable(regs);
10 >         clear_siginfo(&info);
9 >         do_trap(trapnr, signr, str, regs, error_code,
8 >         fill_trap_info(regs, signr, trapnr, &info));
7 >     }
6 }
5
```

结果：

```
lzzz@ubuntu:~/lab07/sig$ dmesg
lzzz@ubuntu:~/lab07/sig$ ./a.out
Hello world!
lzzz@ubuntu:~/lab07/sig$ dmesg
[ 132.839667] Never Giveup
lzzz@ubuntu:~/lab07/sig$
```

4

编写模块task，想法是将已经遍历已经存在的用户进程，对某个用户的进程设置cpuaffinity；然后劫持系统调用sys_execve，对之后建立的进程设置cpuaffinity。

在include/linux/cpumask.h中查看cpumask接口，在kernel/sched/core.c中EXPORT_SYMBOL内核中的sched_setaffinity

```
4822 }
+ 4823 EXPORT_SYMBOL(sched_setaffinity);
4824
```

模块内容如下：

```

3 #include <linux/sched.h>
2 #include <asm/uaccess.h>
1 #include <linux/cpumask.h>
15
1 #define CR0_WP 0x00010000
2 MODULE_LICENSE("GPL");
3 extern void *sys_call_table[];
4
5 static int uid;
6 static int cpunum;
7 module_param(uid, int, 0644);
8 module_param(cpunum, int, 0644);
9
10 unsigned long cr0;
11
12 // extern long sched_setaffinity(pid_t, const struct cpumask *);
13
14 asmlinkage int(*original_call) (const char *, char *const *, char *const *);
15
16
17
18 asmlinkage int my_execve(const char *filename, char *const argv[], char *const envp[])
19 {
20     if (current_uid().val == uid) {
21         struct cpumask cpumask;
22         cpumask_clear(&cpumask);
23         cpumask_set_cpu(cpunum, &cpumask);
24         sched_setaffinity(current->pid, &cpumask);
25     }
26     return original_call(filename, argv, envp);
27 }
28
29
30 int init_module()
31 {
32     struct cpumask cpumask;
33     struct task_struct *ts;
34     cpumask_clear(&cpumask);
35     cpumask_set_cpu(cpunum, &cpumask);
36     for_each_process(ts) {
37         if (task_uid(ts).val == uid)
38             sched_setaffinity(ts->pid, &cpumask);
39     }
40
41     cr0 = read_cr0();
42     write_cr0(cr0 & ~CR0_WP);
43     original_call = sys_call_table[__NR_execve];
44     sys_call_table[__NR_execve] = my_execve;
45     write_cr0(cr0);
46     return 0;
47 }
48
49
50 /*

```

编译之后insmod成功.